

Faculty of Science



Recap: Concurrency Control

Alexander Christensen

Department of Computer Science
University of Copenhagen

2021



Overview

- 1 What is concurrency?
- 2 How can we achieve concurrency?
 - Processes
 - I/O Multiplexing
 - Threading
- 3 What can go wrong in concurrent programming
- 4 Concurrency protection mechanisms




What is concurrency?

"Logical control flows are *concurrent* if they overlap in time." ²

Concurrent Processes

- Each process is a logical control flow.
- Two processes run **concurrently** (are concurrent) if their flows overlap in time
- Otherwise, they are **sequential**
- Examples (running on single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C



² (BOH, cp. 12)



How can we achieve concurrency?

We have 3 different set of tools available to use for writing concurrent programs:

Processes

- A running program can clone itself into one or more child processes
- Completely separated memory areas

I/O Multiplexing

- Resembles event-driven programming
- I/O operations do not block

Threading

- Multiple parallel control flows exist within the same process
- Often the most practical solution



Processes

How do we use them?

Creating Processes

- *Parent process* creates a new running *child process* by calling `fork`
- `int fork(void)`
 - Returns 0 to the child process, child's PID to parent process
 - Child is *almost* identical to parent:
 - Child get an identical (but separate) copy of the parent's virtual address space.
 - Child gets identical copies of the parent's open file descriptors
 - Child has a different PID than the parent
- `fork` is interesting (and often confusing) because it is called *once* but returns *twice*



How to use Processes

Linux interface for using processes: `fork`, `execv`, `getpid`, `wait`, `waitpid`, `exit`.

Concurrency protection mechanisms are not necessary between a process and its child processes because they operate on completely disjoint memory areas.

However, communication between processes is expensive. Common methods are memory mappings (`mmap()`), file pipes (`pipe()`), and sockets.



Processes: Exam question example

Exam question from re-exam 19/20:

```
int main() {  
    int status;  
  
    printf("Hello ");  
    fflush(stdout);  
    printf("%d ", !fork());  
  
    if (wait(&status) != -1)  
        printf("%d ", WEXITSTATUS(status));  
  
    printf("Bye ");  
  
    exit(2);  
}
```

Which of the following are *possible* valid outputs of the program?

- a) Hello 0 1 Bye 2 Bye Yes
- b) Hello Bye 1 0 2 Bye No
- c) Hello 1 0 Bye 2 Bye Yes

- d) Hello 1 Bye 0 2 Bye Yes
- e) Hello 0 Bye 1 2 Bye No
- f) Hello 0 1 Bye Bye 2 No



I/O Multiplexing

I/O operations are blocking. In a webserver `accept()` blocks control flow until it receives an incoming connection request.

But with I/O Multiplexing we can add a number of file descriptors to a read set, e.g. `read-set := {STDIN_FILENO, listenfd}`. This way, a webserver can read from both without blocking, whenever either is ready for reading.

Programming interface: `select`, `FD_ZERO`, `FD_CLR`, `FD_SET`, `FD_ISSET`.

Detail: `select()` is a blocking function.
Example in (BOH, sec. 12.2).



Threading

Using multiple threads (ie. "multi-threading") is the most common way of achieving concurrency in programming.

A Process With Multiple Threads

- **Multiple threads can be associated with a process**
 - Each thread has its own logical control flow
 - Each thread shares the same code, data, and kernel context
 - Each thread has its own stack for local variables
 - but not protected from other threads
 - Each thread has its own thread id (TID)

Thread 1 (main thread)

Thread 2 (peer thread)


Shared code and data

stack 1

stack 2

Thread 1 context:
Data registers
Condition codes
SP1
PC1

Thread 2 context:
Data registers
Condition codes
SP2
PC2



Kernel context:
VM structures
Descriptor table
brk pointer



How to use Threads

Linux (posix) programming interface: `pthread_create`,
`pthread_join`, `pthread_detach`, `pthread_cancel`,
`pthread_exit`, `pthread_self`, `pthread_yield`.

Threads are advantageous in many situations because sharing data btw. peer threads is very easy (`malloc()`), and overhead for creating threads is significantly less than for creating processes. Also, compared to I/O Multiplexing, threads are **truly** parallel.

Threaded programs are exposed to a lot of nasty bugs that can be difficult and time-consuming to track down.



Data sharing in a threaded program

Both global and dynamic (malloc'ed) data is shared between threads:

Example Program to Illustrate Sharing

```
char **ptr; /* global var */

int main()
{
    long i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}
```

sharing.c

```
void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;

    printf("[%ld]: %s (cnt=%d)\n",
        myid, ptr[myid], ++cnt);
    return NULL;
}
```

Peer threads reference main thread's stack indirectly through global ptr variable



What can go wrong in concurrent programming

We will focus only on concurrency issues with multi-threading.

Common issues

- Data races / race conditions
- Deadlocks

Other issues that we omit in this lecture

- Starvation
- Contention / congestion
- Busy-wait loops (extensive resource usage)
- Priority-inversion



Example: What can go wrong (1)

We have a global variable:

```
int* a = malloc(sizeof(int));  
*a = 5;
```

Now, two separate control flows (threads T1 and T2) use this variables concurrently.

T1:

```
int temp = *a;  
temp += 3;  
*a = temp;
```

T2:

```
int temp = *a;  
temp += 4;  
*a = temp;
```

We will call like this:

```
create_new_thread(T1);  
create_new_thread(T2);
```

Will this work?



Example: What can go wrong (2)

```
// int* a = malloc(sizeof(int));  
// *a = 5;
```

In reality, this may occur:

int temp = *a;	(T1)
temp += 3;	(T1)
int temp = *a;	(T2)
*a = temp;	(T1)
temp += 4;	(T2)
*a = temp;	(T2)

What is the result value of *a after both threads have run?

How can we fix this?



Example: What can go wrong (3)

– **FIX: Introducing a mutex:**

T1:

```
lock();  
int temp = *a;  
temp += 3;  
*a = temp;  
unlock();
```

T2:

```
lock();  
int temp = *a;  
temp += 4;  
*a = temp;
```

Will this protect us?



Example: What can go wrong (4)

Oops.. we forget to unlock the mutex in T2. Imagine this control flow:

```
// int* a = malloc(sizeof(int));  
// *a = 5;
```

lock();	(T2)
int temp = *a;	(T2)
temp += 4;	(T2)
*a = temp;	(T2)
lock();	(T1)
int temp = *a;	(T1)
temp += 3;	(T1)
*a = temp;	(T1)
unlock();	(T1)

What will happen?



Example: What can go wrong (5)

More complicated example (a common pitfall !):

```
void modify()  
{  
    lock();  
    int temp = *a;  
    if(temp >= 8) return;  
    temp += 3;  
    *a = temp;  
    unlock();  
}
```

T1:

`modify();`

T2:

`modify();`

Is this correct?



Example: What can go wrong (6)

The problem:

```
void modify()
{
    lock();
    int temp = *a;
    if(temp >= 8) return;
    temp += 3;
    *a = temp;
    unlock();
}
|
```

The solution:

```
void modify()
{
    lock();
    int temp = *a;
    if(temp >= 8) {
        unlock();
        return;
    }
    temp += 3;
    *a = temp;
    unlock();
}
```

Condition variables? See video-recap on A5 on Absalon :)



Concurrency protection mechanisms

Fortunately, we have some powerful tools at our disposal to help us overcome these issues, provided we use them correctly:

Mutex

- Obtain exclusive access to critical sections
- `pthread_mutex_(init|destroy|lock|unlock)`

Condition variable

- Threads go to sleep waiting for a specific condition
- `pthread_cond_(init|destroy|wait|broadcast)`

Semaphore

- Like mutex, except N threads can lock simultaneously
- `sem_init, sem_destroy, sem_wait, sem_post`



Processes: Exam question example

Exam question from re-exam 18/19, covering recap of several OS areas:

<i>For each statement, answer True or False. (Put one "X" in each.)</i>	True	False
a) Before calling <code>thread_cond_wait(cond, mutex)</code> , it is important to unlock <code>mutex</code> first.		X
b) Two processes can share read/write memory via <code>mmap()</code> .	X	
c) To free an array allocated with <code>calloc()</code> , we can pass the address of any element of the array to <code>free()</code> .		X
d) EOF (as returned by e.g. <code>getchar()</code>) is of type <code>unsigned char</code> .		X
e) A semaphore can be implemented with mutexes and condition variables.	X	



Summary

We have seen how to construct concurrent work flows.

We have seen what can go wrong in concurrent programming, and how to fix some common issues.

We have looked at some exam questions from previous years.

Good luck in your exam.



Inspirational quote



C makes it easy to shoot yourself in the foot; C++ makes it harder, but when you do, it blows away your whole leg.

— Bjarne Stroustrup —

AZ QUOTES



Questions

???

